

Security / Crypto Basics

... for C++ Developers

Topics

- Very quick intro
- Getting Random Data
- How to work with passwords?
- How to encrypt data?
- How to communicate securely?

First rule of Crypto: don't do it yourself!

- Getting Cryptography right is REALLY hard
 - Use well tested / reviewed implementations!
- Even using standard libs is hard (cf. OpenSSL)
 - But there are nicer alternatives
- Here are some basic guidelines how NOT TO FAIL immediately.

Types of Algorithms - Hashes

Trap Door Functions / “Hash Functions”

- Simple (non-cryptographic) example: $f(x) = x \% 2$
 - Simple to calculate, impossible to reverse
 - CRC / Checksum / parity: validate data, NOT cryptographic!
- Cryptographic Hashes:
 - very easy to calculate hash for an input
 - very (computationally) difficult to calculate the input (“preimage”) of any given hash.
 - very unlikely that two (even slightly) different messages produce the same value (“collision”)
 - should be “not even close”
- Examples: MD5, SHA-1, SHA-2/3 (256,384,512) 😊
 - (SHA-256 [0..128] is also OK!)

recently broken, don't use (but don't panic ;))

Types of Algorithms - Symmetric

Most secure (in theory): “One Time Pad” - Problem: $\text{len}(\text{key}) == \text{len}(\text{data})$

XOR

Key: 43252987948237957298347598734987598274587...

Data: The quick brown fox jumps over the lazy dog ...

Ciphertext: 624965798629875962557609827967405769...

Block Cipher

Key: 1234	Key: 1234	Key: 1234	Key: 1234	Key: 1234	Key: 1234
The quick	brown fox	jumps over	the lazy	dog <PAD>	The quick
234345234	102340912	234345234

==

Types of Algorithms

Stream Cipher


XOR

Key: 1234 \rightarrow RNG(salt: 1234) \rightarrow 1498759875983789573984759873498572893

Data: The quick brown fox jumps over the lazy dog ...

Ciphertext: 624965798629875962557609827967405769...

Block as Stream Cipher Counter (CTR) or “Chiper Block Chaining” (CBC)



Types of Algorithms - Public Key (RSA or EC)

Private Key

+

The quick brown fox ...

SIG + The quick brown fox ...

+

Public Key

=

Verified

Public Key

+

The quick brown fox ...

Ciphertext: 1349130948108092...

+

Private Key

=

Plaintext

Encryption Strength

Strength (bits)	Hash	Symmetric	Elliptic Curve (~bits)	RSA Modulus / DH Group
56	MD5	DES-56	112	768
80	~SHA-1	2-DES	160	1024
112	SHA2/3-224	3-DES	224	2048
128	SHA2/3-256	AES-128	256	3072
192	SHA2/3-384	AES-192	384	7680
256	SHA2/3-512	AES-256	512	15360

Verordnung Nr. 428/2009 (Dual-Use) 5A002: Symm. 56, EC 112, RSA/DH 512

Crypto Libraries

- “default”: OpenSSL
 - well known / reviewed
 - ugly old API
 - code is a mess, several projects try to fix it up
 - LibreSSL (OpenBSD folks) - close to original API
 - BoringSSL (Google) - more cleanups, diverging API
- Good alternative: Botan
 - nice/safe C++ API
 - SSL/TLS, Crypto Hardware support
 - Recently reviewed / endorsed by BSI
 - 3 year support / bugfix grant
- Crypto++: reasonable C++ API, no SSL

Library Basics: Getting random data

- Good random data essential for secure crypto !
- DON'T use `std::random*` for cryptography
- Use platform library, OpenSSL by default
 - Other options: read from `/dev/random`
 - `CryptGenRandom` on Windows

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

Example: OpenSSL RAND_bytes

Example: Botan - Crypto/SSL for C++

How to work with passwords?

Hashing for Passwords

- Storing/sending “username:password”
 - Eavesdropping / stealing file
- Storing/sending “hash(password)”
 - attacker does not learn user password (if hash / password are secure)
 - but: sending hash(password) \Rightarrow effectively your password now!
 - Use challenge-response (see below)
- Challenge-response:
 - server: “challenge”, client: hash(challenge | pw | nonce), nonce
 - Server needs password (or hash-password)
 - Password/hash stolen from server: attacker can impersonate client

Hashing Problems

- Problems: simple Passwords
 - easy to guess
 - pre-compute often used hashes => Rainbowtable
- Rainbowtables
 - use (good random) SALT: $\text{hash}(\text{salt} + \text{password})$
 - \Rightarrow store “username : SALT : $\text{hash}(\text{salt} + \text{password})$ ”
- Computing $\text{hash}(\text{salt} + \text{password})$ still feasible for known passwords

Better Hashing for Passwords

- Prevent “easy” checking of known passwords: run $\text{HASH}(\text{salt}+\text{PW})$ X times
 - for large X e.g. PBKDF2
- Problem: GPUs still fast enough
 - \Rightarrow BCRYPT: combines HASH with mutating MEM table, kills GPU performance
- Problem: ASICs still fast at BCRYPT
 - \Rightarrow SCRYPT: dynamically growing large MEM table, kills ASIC performance

Example: OpenSSL / Botan Hashing

Storing Passwords

- Never use PLAINTEXT STORAGE!
- Passwd format: \$Algorithm\$Salt\$Hash(Password)

If plaintext PASSWORDS are required (also good for hashes):

- Restrict access to file / verifying process (OS security)
- Separate authentication system
- Encrypt all authentication traffic (e.g also think of NFS, SQL Connections...)
- Encrypt file (key on other medium / TPM) -> safe against discarded disk / stealing file
- Do obfuscate / make access harder, non obvious!

Transmitting Passwords

- Storing hash: protects against storage theft, not eavesdropping!
 - The hash is effectively your new password
- Challenge-response: server stores PW encrypted with e.g. $\text{HASH}(\text{PBKDF2}(\text{PASSWORD}))$, --> access to server store does not reveal hash
- less secure, simple option: just rely on TLS, send plaintext -> same security (-TLS) for you, but exposes user Password !

“State of the art”: “Secure Remote Password” (SRP)

- “DH” for passwords
- Server stores “username: verifier + salt”
 - verifier generated by client, password never transmitted
- Client/server negotiate key using Password(client), Verifier (server)
 - Client → Server: username, random key A
 - Server → Client: salt, random key B
 - If communication works, both generate matching key (may check explicitly or just encrypt)
- Standardized for SSL, available in OpenSSL / Botan
- Only “issue”: stolen verifier can impersonate server
 - So still verify your certs!

How to encrypt data?

Encrypting Data - Use Standard Tools

```
openssl aes-256-ctr -a -salt -in input.txt -out output.txt.enc
```

```
openssl aes-256-ctr -d -a -salt -in input.txt.enc -out output.txt
```

- Problem: unencrypted data written to disk!

Sqlcipher - SQLite with encryption:

- simple encrypted single file data store, secure “by default”

ZIP / LZMA compressors support encryption

Encrypting Data - Standard Algorithm: AES

Use standard algorithm: AES (perhaps Serpent, Blowfish, Stream Cipher: ChaCha20)


- 128, 192 & 256 are regarded as secure
- NOT DES, RC4

What are block chaining modes?

→ turn block cipher into stream, block key based on position

- CBC: key hash based on previous block (don't use)
- ✓ CTR: key based on “block counter” (parallelizable!)
- ✓ GCM: authenticated encryption (otherwise CTR)
 - Additional “tag” validating data
- ✓ XTS: special disk encryption mode

ECB Mode:



(https://en.wikipedia.org/wiki/Block_cipher_mode_of_operation)

Example: OpenSSL / Botan Encryption

Encrypting Data - Disk Encryption

- Block encryption - device level, independent of filesystem
 - No metadata per block → IV based on Sector / Block offset (“tweak”)
 - CBC not well suited: requires reading in sequence, allows “watermarking” attack
 - Solution: XTS mode - encrypts tweak for each block
 - OpenSSL: different keys for tweak / data → 2x key size!
 - Still: no authentication of data! Attacker can revert or destroy blocks

→ Use XTS, better use existing solution (LUKS, VeraCrypt, BitLocker etc)

- Filesystem encryption: e.g. EncFS, NTFS-EFS
 - Can use metadata (e.g. include GCM tag, like ZFS)
 - Often does not encrypt metadata (filename, change times etc)

How to communicate securely?

Asymmetric Encryption: RSA / EC

Public / private key pair:

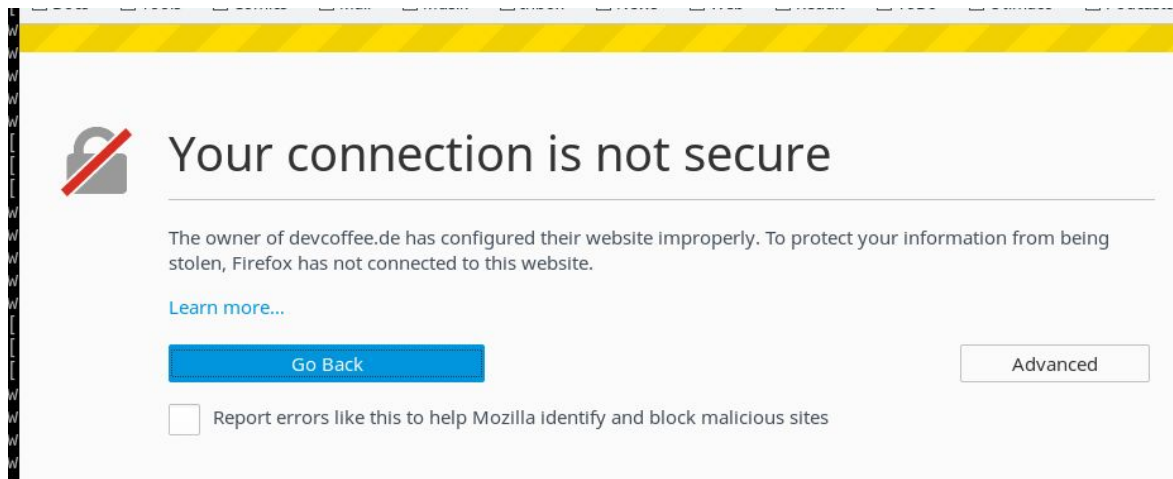
- Private Key encrypts, Public Key decrypts → Signature
- Public Key encrypts, Private Key decrypts → Encryption

Build certificate “chains”: SSL/TLS certificate authorities

- CA certifies intermediate -> intermediate certifies your certificate
- you prove ownership of cert with private key
 - protect your private keys!
 - e.g. again encryption / obfuscation on system

SSL/TLS

- SSL is broken, use TLS 1.2 (1.3)
- if you “Add an exception...”, might as well not encrypt at all:



- validate certificates!
- get CA certified for free: e.g. <https://letsencrypt.org/>

SSL/TLS - Key Negotiation - Diffie-Hellman

- Negotiate secure key without “exchanging” it
- Use “old” DH od ECDH
- Use “ephemeral” mode (DHE/ECDHE)→ “perfect forward secrecy”
- Generate your own DH parameters:

`openssl dhparam -outform pem -out dhparam.pem 2048`

(may take some time, 4096 even more, perhaps pre-compute per release?)

- Problem: man-in-the-middle attacks - relies on CA key verification!

Implementing TLS

- Boost ASIO has good SSL support
 - works around OpenSSL issues
- Botan SSL
- Provide up-to-date CA files or BETTER use OS ones
- Generate (or provide self-generated) DH parameters
- Restrict accepted algorithms
- Implement certificate revocation - way to notify on stolen private keys
 - Note: CRLs not used by browsers, use OSCP

Getting help

If you NEED to DIY, look here:

Books: Practical Cryptography (Ferguson, Schneier), Applied Cryptography (Schneier), Security Engineering (Anderson)


- <https://www.securecoding.cert.org>
- <https://security.stackexchange.com/>

Other Crypto APIs:

- Windows: <https://msdn.microsoft.com/en-us/library/ms867086.aspx>
- Nicer C++ API (buy no SSL) <https://www.cryptopp.com>

Stick Figure Guide to AES: <http://www.moserware.com/2009/09/stick-figure-guide-to-advanced.html>

Happy (and safe) coding!



<https://xkcd.com/538/>