

Security / Crypto Basics

... for C++ Developers

Topics

- Very quick intro
- Getting Random Data
- How to work with passwords?
- How to encrypt data?
- How to communicate securely?

First rule of Crypto: don't do it yourself!

- Getting Cryptography right is REALLY hard
 - Use well tested / reviewed implementations!
- Even using standard libs is hard (cf. OpenSSL)
 - But there are nicer alternatives (e.g. Botan, Crypto++)
- Here are some basic guidelines how NOT TO FAIL immediately.

Types of Algorithms - Hashes

Trap Door Functions / “Hash Functions”

- Simple (non-cryptographic) example: $f(x) = x \% 2$
 - Simple to calculate, hard to reverse
 - CRC / Checksum / parity: validate data, NOT cryptographic!
- Cryptographic Hashes:
 - very easy to calculate hash for an input
 - very difficult to calculate the input (“preimage”) for a given hash.
 - very unlikely that two (even slightly) different messages produce same hash (“collision”)
 - should be “totally unrelated”
- Examples: MD5, SHA-1, SHA-2/3 (256,384,512) 😊
 - (SHA-256 [0..128] is also OK!)

recently broken, don't use (but don't panic ;))

Types of Algorithms - Symmetric

Most secure (in theory): “One Time Pad” - Problem: $\text{len}(\text{key}) == \text{len}(\text{data})$

XOR

Key: 43252987948237957298347598734987598274587...

Data: The quick brown fox jumps over the lazy dog ...

Ciphertext: 624965798629875962557609827967405769...

Block Cipher



Types of Algorithms

Stream Cipher

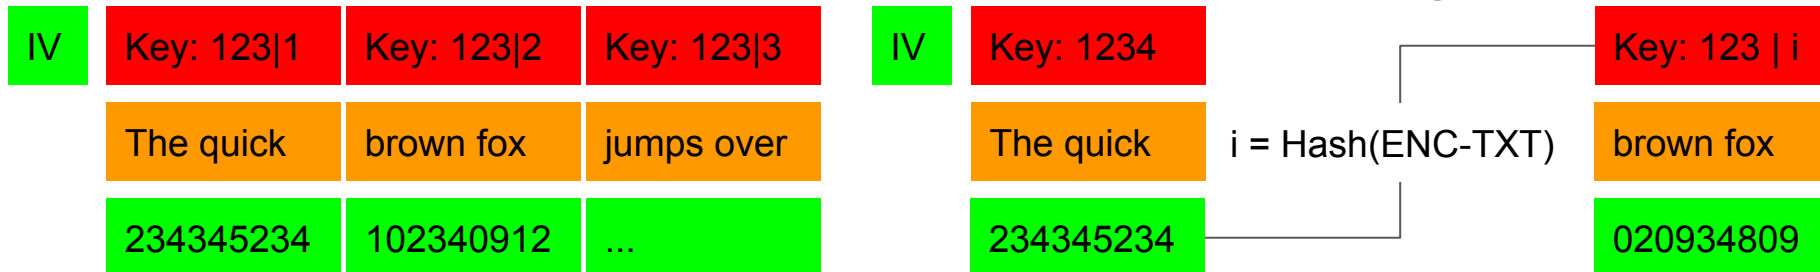
XOR

Key: 1234 \rightarrow RNG(salt: 1234) \rightarrow 1498759875983789573984759873498572893

Data: The quick brown fox jumps over the lazy dog ...

Ciphertext: 624965798629875962557609827967405769...

Block as Stream Cipher Counter (CTR) or “Chiper Block Chaining” (CBC)



Types of Algorithms - Public Key (RSA or EC)



- Slow compared to symmetric
- Usually: encrypt hash (sig) or key (enc) with pub/priv key

Encryption Strength

Strength (bits)	Hash	Symmetric	Elliptic Curve (~bits)	RSA Modulus / DH Group
56	~MD5	DES-56	112	768
80	~SHA-1	2-DES	160	1024
112	SHA 2 / 3 - 224	(3-DES)	224	2048
128	SHA 2 / 3 - 256	AES-128	256	3072
192	SHA 2 / 3 - 384	AES-192	384	7680
256	SHA 2 / 3 - 512	AES-256	512	15360

Verordnung Nr. 428/2009 (Dual-Use) 5A002: Symm. 56, EC 112, RSA/DH 512

Crypto Libraries

- “default”: OpenSSL
 - well known / reviewed
 - ugly old API
 - code is a mess, several projects try to fix it up
 - LibreSSL (OpenBSD folks) - close to original API
 - BoringSSL (Google) - more cleanups, diverging API
- Good alternative: Botan
 - nice/safe C++ API
 - SSL/TLS, Crypto Hardware support
 - Recently reviewed / endorsed by BSI
 - 3 year support / bugfix grant
- Crypto++: reasonable C++ API, no SSL

Library Basics: Getting random data

- Good random data essential for secure crypto !
- DON'T use `std::random*` for cryptography
- Use platform library, OpenSSL by default
 - Other options: read from `/dev/random`
 - `CryptGenRandom` on Windows

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

Example: OpenSSL RAND_bytes

Example: Botan - Crypto/SSL for C++

How to work with passwords?

Hashing for Passwords

- Storing/sending “username:password”
 - Eavesdropping / stealing file
- Storing/sending “hash(password)”
 - attacker does not learn user password (if hash / password (salt) are secure)
 - but: hash(password) effectively the password now!
- Challenge-response:
 - server: “challenge”, client: hash(challenge | pw | nonce), nonce
 - Server needs password (or hash-password) !
 - Password/hash stolen from server: attacker can impersonate client

Hashing Problems

- Problems: simple Passwords
 - easy to guess
 - pre-compute often used hashes => Rainbowtable
- Rainbowtables
 - use (good random) SALT: $\text{hash}(\text{salt} + \text{password})$
 - bad salt: username (cf. databases - rainbow tables for postgres, dba, sa ...)
- Computing $\text{hash}(\text{salt} + \text{password})$ still feasible for known passwords
 - Password lists or leaked unsalted passwords
 - People reuse their passwords!

Better Hashing for Passwords

- Prevent “easy” checking of known passwords: run $\text{HASH}(\text{salt}+\text{PW})$ X times
 - for large X e.g. PBKDF2
 - iteration count → restrict hashes / sec, measure
 - don't overload your own systems
- Problem: GPUs still fast enough
 - ⇒ BCrypt: combines HASH with mutating MEM table, kills GPU performance
- Problem: ASICs still fast at BCrypt
 - ⇒ SCrypt: dynamically growing large MEM table, kills ASIC performance

Example: OpenSSL / Botan Hashing

Storing Passwords

- Never use PLAINTEXT STORAGE!
- Passwd format: Username:\$Algorithm\$Salt\$Hash(Password)

If plaintext PASSWORDS are required (also good for hashes):

- Restrict access to file / verifying process (OS security)
- Separate authentication system
- Encrypt all authentication traffic (e.g also think of NFS, SQL Connections...)
- Encrypt file (safe key / TPM) → protect against discarded disk / stealing file
- Do obfuscate / make access harder, non obvious!

Transmitting Passwords

- Storing/sending hash: protects against storage theft, not eavesdropping!
 - The hash is effectively your new password
 - But User password will not be revealed by breach
- “Double Hash”: store $\text{Salt:PBKDF2}(n, \text{PBKDF2}(n, \text{Salt}|\text{Password}))$
 - client computes first hash, server the second
 - access to server store does not reveal client hash
 - access to communication still shows client “password”
- Less secure, simple option: just rely on TLS, send plaintext
 - Many services actually do this
 - Certificate validation very important!

“State of the art”: “Secure Remote Password” (SRP)

- “DH” for passwords
- Server stores “username: verifier + salt”
 - verifier generated by client, password never transmitted
- Client/server negotiate key using Password(client), Verifier (server)
 - Client → Server: username, random key A
 - Server → Client: salt, random key B
 - If communication works, both generate matching key (may check explicitly or just encrypt)
- Standardized for SSL, available in OpenSSL / Botan
- Only “issue”: stolen verifier can impersonate server
 - So still verify your certs!

How to encrypt data?

Encrypting Data - Use Standard Tools

```
openssl aes-256-ctr -a -salt -in input.txt -out output.txt.enc
```

```
openssl aes-256-ctr -d -a -salt -in input.txt.enc -out output.txt
```

- Problem: unencrypted data written to disk!

Sqlcipher - SQLite with encryption:

- simple encrypted single file data store, secure “by default”

ZIP / LZMA compressors support encryption

Encrypting Data - Standard Algorithm: AES

Use standard algorithm: AES (perhaps Serpent, Blowfish, Stream Cipher: ChaCha20)


- 128, 192 & 256 are regarded as secure, additional parameter: rounds (usually fixed)
- NOT DES, RC4

What are block chaining modes? → turn block cipher into stream

- CBC: key hash based on previous block + IV
 - Note: IV randomness is important (also for GCM)!
- ✓ CTR: key based on “block counter” (parallelizable!)
- ✓ GCM: authenticated encryption (otherwise CTR)
 - Additional “tag” validating data
- ✓ XTS: special disk encryption mode

!!! CTR/GCM: “same IV same key” problem: reveal XOR(PT1, PT2)

ECB Mode:



Example: OpenSSL / Botan Encryption

Encrypting Data - Disk Encryption

- Block encryption - device level, independent of filesystem
 - Several solutions: LUKS, True/VeraCrypt, BitLocker etc.
 - No metadata per block → IV based on Sector / Block offset (“tweak”)
 - CBC problem: requires reading in sequence, allows “watermarking” attack
 - Solution: **XTS** mode - encrypts tweak for each block
 - OpenSSL: different keys for tweak / data → 2x key size!
 - Still: no authentication of data! Attacker can revert or destroy blocks
- Filesystem encryption: e.g. EncFS, NTFS-EFS
 - Can use metadata (e.g. include GCM tag, like ZFS, EncFS currently does NOT!)
 - Often does not encrypt file metadata (filename, change times etc)

How to communicate securely?


Asymmetric Encryption: RSA / EC

RSA: “factoring large primes”

- Primes $p, q \rightarrow$ Modulus $n = p \cdot q$, public/private exponents e, d
 - $\rightarrow (m^e)^d \equiv m \pmod{n}$
- Private key: (d, n) , public key: (e, n)
- Security: “modulus size”

Elliptic Curves - “elliptic curve *discrete* logarithm”

- Choose base curve (e.g. SECP256,384,521, Curve25519) \rightarrow security level
- Public key: point $A = n \cdot G$ (generator), private key: number of “hops” n
 - Encrypt: $R = r \cdot G$, $S = r \cdot A$ send $R \rightarrow S = n \cdot R$
- Signature: ECDSA, Key Exchange: ECDHE, Encryption: ECIES



https://en.wikipedia.org/wiki/Elliptic_curve

<https://arstechnica.com/security/2013/10/a-relatively-easy-to-understand-primer-on-elliptic-curve-cryptography/>

SSL/TLS

Build certificate “chains”: SSL/TLS certificate authorities

- CA certifies intermediate → intermediate certifies domain certificate
- you prove ownership of cert with private key
 - protect your private keys!
 - e.g. again encryption / obfuscation on system
- Validation: $\text{hash}(\text{cert}) \rightarrow$ security depends on hash!
 - MD5 certs are insecure, SHA1 may be
- Stolen private keys: certificate revocation!
 - using CRLs: clients periodically update revocation lists (Browsers don't !!!)
 - using OSCP: active query for each connection (Browsers: “soft fail” or nothing (Chrome))

SSL/TLS - Key Negotiation - Diffie-Hellman

- Negotiate secure key without “exchanging” it
- Use “old” DH or ECDH
- Use “ephemeral” mode (DHE/ECDHE)→ “perfect forward secrecy”
- Generate your own DH parameters:

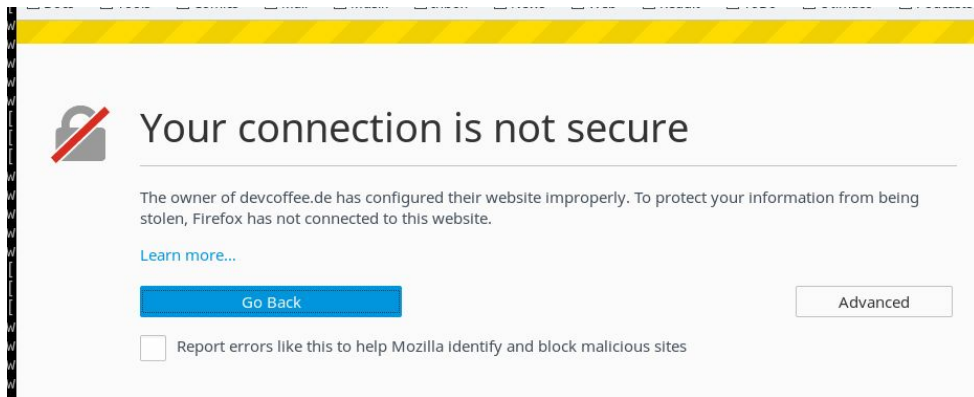
```
openssl dhparam -outform pem -out dhparam.pem 2048
```

(may take some time, 4096 even more, perhaps pre-compute per release?)

- Problem: man-in-the-middle attacks - relies on CA key verification!

SSL/TLS

- SSL is broken, use TLS 1.2 (1.3)
- if you “Add an exception...”, might as well not encrypt at all:



- Validate server certificates!
- Get CA certified (for free: e.g. <https://letsencrypt.org/>)

Implementing TLS

- Boost ASIO has good SSL support
 - works around OpenSSL issues
- Botan SSL
 - with ASIO or sockets as “transport”
- Provide up-to-date CA files or BETTER use OS ones
 - Think about revocation, make sure it works (especially offline!)
- Generate (or provide self-generated) DH parameters
- Restrict accepted algorithms

Getting help

If you NEED to DIY, look here:

Books: Practical Cryptography (Ferguson, Schneier), Applied Cryptography (Schneier), Security Engineering (Anderson)

- <https://www.securecoding.cert.org>
- <https://security.stackexchange.com/>

Other Crypto APIs:

- Windows: <https://msdn.microsoft.com/en-us/library/ms867086.aspx>
- Nicer C++ API (buy no SSL) <https://www.cryptopp.com>

Stick Figure Guide to AES: <http://www.moserware.com/2009/09/stick-figure-guide-to-advanced.html>

Happy (and safe) coding!

